

FACHHOCHSCHULE LUZERN HSLU

STUDIENGANG DIGITAL IDEATION, BACHELOR
4. SEMESTER

Procedural Generation in the age of AI

Simon Hischier

May 21, 2018

Inhalt

1	Abstract	2
2	Definitions	2
2.1	Procedural Content Generation	2
2.2	Artificial Intelligence	2
2.3	The Difference between AI, Machinelearning, Deeplearning	3
2.4	Often used Neural Networks	4
2.4.1	Convolution Neural Network	4
2.4.2	Recurrent Neural Networks	4
2.4.3	Long short-term memory	4
3	The artistic vision and the generation	5
4	Procedural Content Generation and AI	5
5	Applications of Procedural Approaches	6
5.1	Level Generation	6
5.2	Text to Speech	6
6	Applications of AI systems	6
6.1	Face representation in games	6
6.2	evolution	7
6.3	game mechanics	7
6.4	art etc.	7
6.5	text to speech	7
6.6	Game AI	7
6.7	Anti-Cheating	8
6.8	Level Generation compared	8
7	shortcomings of AI	8
8	Procedural Content Generation and Machine Learning	8
9	Conclusion	8
10	Referenzen und Akronyme	8

1 Abstract

blibliablu todo

Fragestellung: How does content creation in games differ between [Procedural Content Generation \(PCG\)](#) and [Machine Learning \(ML\)](#) content creation?

2 Definitions

2.1 Procedural Content Generation

The term [PCG](#) defines a form of content generation that is automated. The content is created through algorithmic processes and with few to no human interaction. This allows for the creation of bigger, more random looking, unique content with less to no artist interaction.[1]

2.2 Artificial Intelligence

In 1956 the Book Automata Studies[2] laid the ground work for [Artificial Intelligence \(AI\)](#) and the Dartmouth summer research project on artificial intelligence marked the key event "to nail the flag to the mast." McCarthy is credited for coining the phrase "artificial intelligence" and solidifying the orientation of the field[3]. The name [AI](#) was used ever since for various applications. Ever since this key event [AI](#) is defined based on the goal that is tried to beeing achieved. Bernard Marr lists them as following:[4]

1. Build systems that think exactly like humans do ("strong AI")
2. Just get systems to work without figuring out how human reasoning works ("weak AI")
3. Use human reasoning as a model but not necessarily the end goal

Marr referes with "strong AI" and "weak AI" to the paper written by John Searle where he defines a strong [AI](#) of beeing able to think and have a mind and a weak [AI](#) that can only act like it thinks and has a mind. The paper is also known for the "Chinese Room" argument[5]. More terms for a [AI](#) classification exist like Artificial General Intelligence (AGI) and the Artificial Superintelligence (ASI). These classifications are not further defined in this work.

Bernard Marr continues to list the various definitions for [AI](#). While the dictionaries list [AI](#) as a definition, companies lack a clear definition and Marr extrapolates a definition from the companies research field. The dictionary definitions and his extrapolated definitions are listed here:

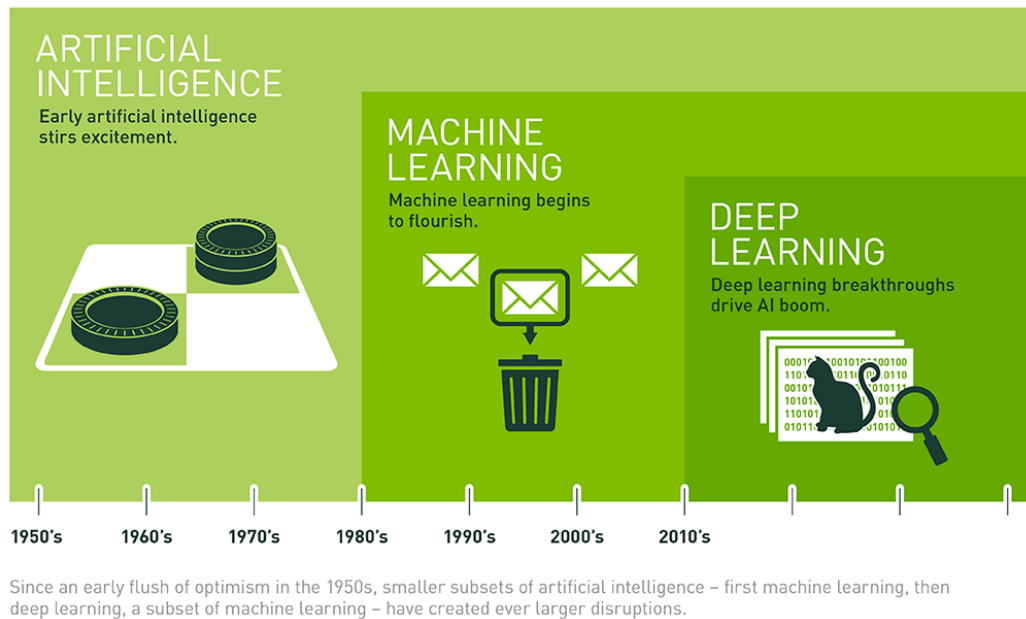
1. **The English Oxford Living Dictionary** "The theory and development of computer systems able to perform tasks normally requiring human intelligence, such as visual perception, speech recognition, decision making, and translation between languages."
2. **Merriam-Webster**
 - (a) A branch of computer science dealing with the simulation of intelligent behavior in computers.
 - (b) The capability of a machine to imitate intelligent human behavior.
3. **The Encyclopedia Britannica** "artificial intelligence (AI), the ability of a digital computer or computer-controlled robot to perform tasks commonly associated with intelligent beings."

4. **Amazon** defines it as "the field of computer science dedicated to solving cognitive problems commonly associated with human intelligence, such as learning, problem solving, and pattern recognition."
5. **Google AI** "create smarter, more useful technology and help as many people as possible"
6. **Facebook AI Research** "advancing the field of machine intelligence and are creating new technologies to give people better ways to communicate."
7. **IBM's** three areas of focus are "AI Engineering, building scalable AI models and tools; AI Tech where the core capabilities of AI such as natural language processing, speech and image recognition and reasoning are explored and AI Science, where expanding the frontiers of AI is the focus." [4]

This work is using the definition of the term **AI** as listed in *The Encyclopedia Britannica*. The fields described in this work can be categorized in the 3. objective listed by Marr "Use human reasoning as a model but not necessarily the end goal" [4]

2.3 The Difference between AI, Machinelearning, Deeplearning

As shown before, **AI** is not a new phenomena and even **ML** is known for quite some time. There is not a great difference between **AI**, **ML** and deep learning. They are subsets of the previous. While **AI** still is a very broad term **ML** and deep learning are closer together. Deep learning introduces neuron simulation to **ML** by creating discrete layers, connections and directions of data propagation[6]. The technology started to grow significantly around the year 2010. This seems to be related to the growing computational power and the immense amount of collected and stored data[6]. The illustration is showing the relationship between the 3 terms.



§

Figure 1: Levels of AI as Image[6].

2.4 Often used Neural Networks

2.4.1 Convolution Neural Network

The [Convolution Neural Network \(CNN\)](#) describes a method to reduce input data and enhance [Neural Network \(NN\)](#) performance. The convolution allows for faster calculations and reduces memory requirements. This advantages make [CNN](#) very desirable for image recognition and voice recognition tasks and are widely used in these fields.

2.4.2 Recurrent Neural Networks

As explained in [subsection 2.3](#) [ML](#) is a broad term and includes a variety of models. The [Recurrent Neural Networks \(RNN\)](#) are networks for tasks where we need some kind of persistence. If we want to classify video frames the network should have some kind of consistency.[7] A network should persist the last seen data and not reclassify items every frame. Reclassifying without previous context could lead to different recognitions in every frame for the same object.

2.4.3 Long short-term memory

[RNN](#) are good for persisting very recent information. Sentences are a great example: "Ships are built to float on *water*". The [RNN](#) is great in filling the end of this sentence. Problems arise when the information is needed a lot later. The more information is inbetween the contextual references the more unreliable a basic [RNN](#) gets. Books for example can have references on the

last page to the very beginning. For such tasks a [Long short-term memory \(LSTM\)](#) model is the perfect fit. The [LSTM](#) network was introduced by Hochreiter Sepp and Uergen Schmidhuber[8]. A [LSTM](#) is a specialized version of a [RNN](#) which is designed for these kind of tasks. Almost all [RNN](#) tasks can be achieved with a [LSTM RNN](#)[7].

3 The artistic vision and the generation

Generating content for games is a fundamental artistic choice for game developers. The generation in various forms is linked to a decrease in the artistic vision. Designers have to step away from micro-controlling game parts like environment, shapes, colors, enemy behaviours etc. Therefore games do include [PCG](#) in different ways and in various depths. Big studios tend to stick to more controlled experiences and have more (human-)resources to ensure this vision. We define a list of various depths of [PCG](#):

1. **No generation** (ex. *Super Mario Bros* (Nintendo, 1985) where everything was handplaced, drawn and animated as explained in *Super Mario Bros*. Level 1-1[9])
2. **Content generation in the game making phase.** (ex. *The Elder Scrolls Oblivion* (Bethesda Softworks, 2006) used [PCG](#) to generated most of the world before the artists curated it.[10] An example of a widely used [PCG](#) algorithm middleware for game studios is SpeedTree[11])
3. **Gameplay (partially) definded or influenced** by [PCG](#) such as the sidequests for *The Elder Scrolls Skyrim* (Bethesda Softworks, 2011) which were endlessly generated[12] or Castles in *Rogue Legacy* (Cellar Door Games, 2013) which are generated procedurally but the game has some kind of continuosity and progress on top of the castle runs.[13]
4. **Games almost completely generated** ex. *Dwarf Fortress* (Dwarf Fortress, 2006) doesn't stop at the map generation. It starts out generating the history of this world and everything that happened before.[14]

Games and even game genres do fall into these different levels of [PCG](#). A major role for this classification of games and game genres is the depth of artistic control or lack therefore.[1] Games that do rely more on [PCG](#) tend to focus more on the fun gameplay rather than an intriguing story and complex characters. Further categorization and different methods of [PCG](#) can be found in the Paper from van der Linden, Lopes and Bidarra[1].

4 Procedural Content Generation and AI

As explained in [The artistic vision and the generation](#) there are various genres and games falling into different levels of [PCG](#). Due to the type of gameplay and game mechanics linked to [PCG](#), games with high levels of procedural generation can easily be identified. We hope [AI](#) blurs the line between [PCG](#) and handcrafted more. A goal for [AI](#) components is to create [PCG](#) games that are less distinguishable from handcrafted counterparts. By extending the [PCG](#) with [AI](#) we hope to have a more natural and handcrafted feel to [PCG](#) type games. As [ML](#) gets used more in games, the applications starts to vary further.

5 Applications of Procedural Approaches

Procedural generation of content varies wildly and can affect most gameparts. In [The artistic vision and the generation](#) we made an attempt to categorize the different stages and depths of automation. However the reason to use [PCG](#) can range from necessity (memory limitations, more content to create then time to create it) to artistic choice (pattern and style generation). The [PCG](#) algorithms allow to shift the artists focus away from the detail work to broader concepts like world biomes in *Minecraft* (Markus Persson, 2009). [PCG](#) has multiple facettes and listing the different reasons to use [PCG](#) is not feasible. What all procedural algorithms share in common is a tradeoff between disk space and computing time to be chosen. A big advantage of [PCG](#) is the great influence on the outcome and the artists controll over what exactly gets generated.

5.1 Level Generation

Procedural level generation is around for a very long time. It was present during the arcade games because of the need to possibly play endlessly to get the high score. It got carried over to the personal computer, generating dungeons to provide unique experiences[1] and to counter floppy disk memory limitations. Currently one of the most known level generators is used in *Minecraft* (Markus Persson, 2009) where the endless worlds are used to feed the players curiosity. The strength of these systems is that worlds are replicable (see world-seeds in *Minecraft*), can be controlled and fine-tuned and don't necessarily need a lot of computational power. They are expandable, maintainable and errors in the level generation process can be checked and fixed. A drawback is the complexity (or lack therefore) of such generators. Worlds can look repetitive and empty.

5.2 Text to Speech

Procedural approaches of text to speech up until now were more on the functional side. Text to speech has great challenges and voices can sound very monotone, lack highlights and are in general not used in games. Games either hire voice actors to read the text examples or display the text as is.

6 Applications of AI systems

TO DO

6.1 Face representation in games

An area where [AI](#) is much more effective then handwritten procedrual approaches is face reconstruction and face mapping. Webcam feeds have no depth information and creating 3D faces from photos is a laborintensive work. Actual research shows [ML](#) is capable of reconstruction and position maping 3D avatars from single images. It's robust, fast and stable[15]. A large area in the games industry is avatar faces. Massive Multiplayer Games tend to have elaborate avatar creaton tools to customize the game character. Games like *Arma* (Bohemia Interactive, 2013) are using the users voice in multiplayer sessions and try to synchronize the avatars lips with the players spoken words. Online chats with avatar representations such as *VRChat* (VRChat Inc., 2017) have immersive VR worlds where players can walk and talk in a virtual environment. For streaming platforms like *Twitch* (Amazon, 2011) a webcam feed is part of the majority of

streams. All these different games, genres and platforms would benefit from some kind of reconstruction of the players face. (Arma) could have the real users faces represented in real-time in the game, non human avatars could use [AI](#) to map the users lips and expressions on to the virtual avatars. Streamer could have their faces as 3D models or have an avatar instead of a live camera feed.

6.2 evolution

asdf

6.3 game mechanics

ai game mechanics?

6.4 art etc.

ex. color palettes style transfer

6.5 text to speech

Text is a big part of games. Extreme examples of games with word counts comparable to The Lord of the Rings book trilogy are present in the [Role Play Game \(RPG\)](#) and text adventure genre. For this games voicing every line is out of range for game studios and therefore players have to read most of the text. With a trained [NN](#) it is possible to have a voice actor read example texts and let the [NN](#) generate the voice from the written text in the game[16]. Although a demonstration used 24.6 hours of training data this number can probably be reduced. Another paper shows, that even singing can be generated with 16 and 35 minutes of sentences read out[17]. The later audio example lengths could potentially open doors for further applications of text to speech such as online chat using the users voice to generate spoken words for the other users.

[18] painting with extra detail

6.6 Game AI

Something almost all games have is computer players controlled by [AI](#). There are a variety of different solutions to mimic human behaviour. The more complex a game is the harder it is to write comprehensible game [AI](#). Games like Go remain too hard to calculate a winning strategy. There are simply too many complex situations and too many possible paths to explore. In this field [AI](#) really showed its power and Alpha Go defeating the world champion was a major step in [AI](#)[19][20]. While Alpha Go did need more hardwarepower then the usual gaming computer provides, OpenAI demonstrated a less powerhungry [AI](#) for *Dota 2* (Valve Corporation, 2013). At the International Tournament 2017 OpenAI demonstrated an [AI](#) that beat Pro Players at the game reliably.[21] The rules were more strict then a normal game of Dota 2 but given enough time and resources, an [AI](#) that can beat Professional teams is likely. These types of games are a big challenge to program computer players for because they allow for wacky situations, situations never seen before and asymmetric or incomplete information about the other players. A strong field for [ML](#) is creating [AI](#) that always remains and adapts to the level of the human player. This is almost impossible to program by hand or with other approaches.

6.7 Anti-Cheating

Valve demonstrated a very significant increase in Anti-Cheat software effectiveness after they started to use deep learning. Old methods of hardcoding different checks did and does only lead to an arms race with the cheat providers. The new approach requires the company to maintain a powerful server. A big drawback is, that the company has to own a more powerful NN and have more complete information to not get outsmarted by cheat providers[22].

6.8 Level Generation compared

Level generation can be done via AI or PCG depending on the desired output. Both approaches could potentially be combined for a wider variety. What both approaches lack is the carefully curated and manually designed human aspect of levels. The amount of detail and teaching in a single *Super Mario Bros.* (Nintendo, 1985) level[9] is still not feasible by both approaches.

7 shortcomings of AI

TO DO , generation RPGs started to rely more on generation and artists curation of content to fill the initial world. -even more lack of control, very big input array -same problems as with PCG because a more complex understanding is needed for an artistic vision across levels

8 Procedural Content Generation and Machine Learning

TO DO The different strengths of PCG and ML suggest, that ML is an extension in the content generation toolbox rather than a replacement.

super mario level 1-1 teaching ...

AI is a new type of content generation with overlapping areas to PCG and classic game AIs. It can go from a PCG replacement to games revolving around an AI (Alien game colonial marines?)

9 Conclusion

? while pcg and ai both have very real and usefull ... they still lack the thought through creative possibilities possible by humans. showed in the example of level generation how the different approaches are ai is harder to controll then PCG and depending on the area is a lot more labor intensive. But has fields where PCG is inappropriate. ai needs more calculation power, big initial work know what to do upfront ?

10 Referenzen und Akronyme

Acronyms

AI Artificial Intelligence. 2, 3, 5–8

CNN Convolution Neural Network. 4

LSTM Long short-term memory. 5

ML Machine Learning. [2–8](#)

NN Neural Network. [4, 7, 8](#)

PCG Procedural Content Generation. [2, 5, 6, 8](#)

RNN Recurrent Neural Networks. [4, 5](#)

RPG Role Play Game. [7, 8](#)

References

- [1] Roland Van Der Linden, Ricardo Lopes, and Rafael Bidarra. Procedural generation of dungeons. *IEEE Transactions on Computational Intelligence and AI in Games*, 6(1):78–89, mar 2014.
- [2] John McCarthy and Claude Elwood Shannon. *Automata Studies: Annals of Mathematics Studies. Number 34 - William Ross Ashby*. Princeton University Press, 1956.
- [3] James Moor. The Dartmouth College artificial intelligence conference: The next fifty years. *Ai Magazine*, 27(4):87, 2006.
- [4] Bernard Marr. The Key Definitions Of Artificial Intelligence (AI) That Explain Its Importance, 2018.
- [5] John R. Searle. Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(03):417, sep 1980.
- [6] Michael Copeland(Nvidia). The Difference Between AI, Machine Learning, and Deep Learning? — NVIDIA Blog, 2016.
- [7] Nery Sofiyanti, Dyah Iriani Fitmawati, and Andesba A. Roza. Understanding LSTM Networks, 2015.
- [8] Sepp Hochreiter and Jürgen Schmidhuber. Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780, 1997.
- [9] Eurogamer. Miyamoto on World 1-1, 2015.
- [10] Bethesda Game Studios. The Elder Scrolls 5: Skyrim, 2011.
- [11] SpeedTree. SpeedTree Vegetation Modeling.
- [12] Matt Bertz. The Technology Behind The Elder Scrolls V: Skyrim, 2011.
- [13] Rich Stanton. The making of Rogue Legacy • Eurogamer.net, 2013.
- [14] Alex J. Champanand. Beyond Terrain Generation: Bringing Worlds in DWARF FORTRESS to Life — AiGameDev.com, 2012.
- [15] Yao Feng, Fan Wu, Xiaohu Shao, Yanfeng Wang, and Xi Zhou. Joint 3D Face Reconstruction and Dense Alignment with Position Map Regression Network. mar 2018.

- [16] Jonathan Shen, Ruoming Pang, Ron J. Weiss, Mike Schuster, Navdeep Jaitly, Zongheng Yang, Zhifeng Chen, Yu Zhang, Yuxuan Wang, RJ Skerry-Ryan, Rif A. Saurous, Yannis Agiomyrgiannakis, and Yonghui Wu. Natural TTS Synthesis by Conditioning WaveNet on Mel Spectrogram Predictions. dec 2017.
- [17] Merlijn Blaauw and Jordi Bonada. A Neural Parametric Singing Synthesizer. apr 2017.
- [18] Fujun Luan, Sylvain Paris, Eli Shechtman, and Kavita Bala. Deep Painterly Harmonization. apr 2018.
- [19] David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, jan 2016.
- [20] David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of Go without human knowledge. *Nature*, 550(7676):354–359, oct 2017.
- [21] Openai. More on Dota 2, 2017.
- [22] John McDonald. GDC Vault - Robocalypse Now: Using Deep Learning to Combat Cheating in 'Counter-Strike: Global Offensive', 2018.

List of Figures

1	Levels of AI as Image[6].	4
---	-----------------------------------	---

List of Tables